

Machine Learning

August 6, 2023

1 Machine Learning

```
[1]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
```

```
[2]: %matplotlib inline
```

1.1 Loading Data

```
[3]: ecg_data = pd.read_csv("./processed_data/ecg_processed_data.csv", index_col=0)
```

```
[4]: ecg_data.head()
```

```
[4]:
```

	Record ID	Segment	Start	Segment End	0	1	2	\
0	100		13	213	0.059449	0.055308	0.052035	
1	100		307	507	0.084239	0.087566	0.088956	
2	100		563	763	-0.006937	-0.005690	-0.005184	
3	100		883	1083	0.072086	0.073128	0.073026	
4	100		1168	1368	0.084762	0.083995	0.082327	

	3	4	5	6	...	191	192	193	\
0	0.049666	0.047961	0.046922	0.046801	...	0.029812	0.031902	0.033233	
1	0.088621	0.086919	0.084459	0.081823	...	0.020380	0.020170	0.020049	
2	-0.005172	-0.005405	-0.005802	-0.006385	...	-0.083923	-0.084476	-0.084798	
3	0.071802	0.069731	0.067336	0.065156	...	0.040917	0.040820	0.040605	
4	0.080127	0.077701	0.075216	0.072963	...	0.024691	0.025225	0.025977	

	194	195	196	197	198	199	\
0	0.033981	0.034276	0.034166	0.033702	0.033134	0.032648	
1	0.020272	0.020719	0.020971	0.020866	0.020622	0.020434	
2	-0.084841	-0.084325	-0.082953	-0.080644	-0.077424	-0.073181	
3	0.040177	0.039328	0.038051	0.036671	0.035565	0.034817	
4	0.027249	0.029211	0.031715	0.034364	0.036462	0.037443	

Annotation Class

```

0          N
1          N
2          N
3          N
4          N

```

[5 rows x 204 columns]

```
[5]: ecg_data["Annotation Class"].unique()
```

```
[5]: array(['N', 'V', '/', 'L', 'R'], dtype=object)
```

```
[6]: ecg_data["Annotation Class"].nunique()
```

```
[6]: 5
```

```
[7]: ecg_data["Annotation Class"].value_counts()
```

```

[7]: Annotation Class
N      73439
L       8068
R       7255
V       6793
/       3619
Name: count, dtype: int64

```

1.2 Preprocessing

1.2.1 Converting Non Numeric Column to Numeric discontinuous Columns

```

[8]: annotation_dict = dict()
for i, symbol in enumerate(ecg_data["Annotation Class"].unique()):
    annotation_dict[symbol] = i + 1
ecg_data["Annotation Class Numeric"] = ecg_data["Annotation Class"].apply(
    lambda x: annotation_dict[x]
)

```

```
[9]: ecg_data.head()
```

```

[9]:   Record ID  Segment Start  Segment End      0      1      2 \
0         100           13        213  0.059449  0.055308  0.052035
1         100          307        507  0.084239  0.087566  0.088956
2         100          563        763 -0.006937 -0.005690 -0.005184
3         100          883       1083  0.072086  0.073128  0.073026
4         100         1168       1368  0.084762  0.083995  0.082327

      3      4      5      6 ...      192      193      194 \

```

```

0  0.049666  0.047961  0.046922  0.046801  ...  0.031902  0.033233  0.033981
1  0.088621  0.086919  0.084459  0.081823  ...  0.020170  0.020049  0.020272
2 -0.005172 -0.005405 -0.005802 -0.006385  ... -0.084476 -0.084798 -0.084841
3  0.071802  0.069731  0.067336  0.065156  ...  0.040820  0.040605  0.040177
4  0.080127  0.077701  0.075216  0.072963  ...  0.025225  0.025977  0.027249

```

```

          195          196          197          198          199  Annotation Class \
0  0.034276  0.034166  0.033702  0.033134  0.032648                      N
1  0.020719  0.020971  0.020866  0.020622  0.020434                      N
2 -0.084325 -0.082953 -0.080644 -0.077424 -0.073181                      N
3  0.039328  0.038051  0.036671  0.035565  0.034817                      N
4  0.029211  0.031715  0.034364  0.036462  0.037443                      N

```

```

Annotation Class Numeric
0                      1
1                      1
2                      1
3                      1
4                      1

```


[5 rows x 205 columns]


```
[10]: x_columns = [str(i) for i in range(0, 200)]
```


1.3 What each Annotation looks like?


```
[11]: for symbol in ecg_data["Annotation Class"].unique():
        readings = (
            ecg_data[ecg_data["Annotation Class"] == symbol].head(1)[x_columns].
            values[0]
        )
        plt.figure(figsize=(10, 8))
        plt.title(label=symbol)
        plt.plot(readings)
        plt.show()


```











1.4 Splitting into Test and Train Data

```
[12]: from sklearn.preprocessing import MinMaxScaler
```

```
[13]: from sklearn.pipeline import Pipeline
```

```
[14]: from sklearn.model_selection import train_test_split
```

```
[15]: X = ecg_data[x_columns].values
```

```
[16]: y = ecg_data["Annotation Class Numeric"].values
```

```
[17]: X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=101, stratify=y, shuffle=True
)
X_train_1, X_val, y_train_1, y_val = train_test_split(
    X_train, y_train, test_size=0.2, stratify=y_train, shuffle=True,
    random_state=101
```



```
)
```

```
[18]: X_train.shape
```

```
[18]: (79339, 200)
```

```
[19]: y_train.shape
```

```
[19]: (79339,)
```

```
[20]: X_test.shape
```

```
[20]: (19835, 200)
```

```
[21]: y_test.shape
```

```
[21]: (19835,)
```

1.5 Logistic Regression

```
[22]: from sklearn.linear_model import LogisticRegression  
      from sklearn.model_selection import GridSearchCV
```

```
[23]: from skopt import BayesSearchCV
```

```
[24]: lr_params = {"multi_class": ["ovr", "multinomial", "auto"]}
```

```
[25]: pipeline = Pipeline(  
    [  
        ("minmax", MinMaxScaler()),  
        ("lr", GridSearchCV(LogisticRegression(), lr_params, verbose=3)),  
    ]  
)
```

```
[26]: pipeline.fit(X_train, y_train)
```

Fitting 5 folds for each of 3 candidates, totalling 15 fits

```
/home/alton/.local/lib/python3.9/site-  
packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-

```
regression
  n_iter_i = _check_optimize_result(
[CV 1/5] END ..multi_class=ovr;, score=0.839 total time= 11.4s
/home/alton/.local/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear_model.html#logistic-](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

regression

```
  n_iter_i = _check_optimize_result(
[CV 2/5] END ..multi_class=ovr;, score=0.843 total time= 13.2s
/home/alton/.local/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear_model.html#logistic-](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

regression

```
  n_iter_i = _check_optimize_result(
[CV 3/5] END ..multi_class=ovr;, score=0.843 total time= 10.8s
/home/alton/.local/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear_model.html#logistic-](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

regression

```
  n_iter_i = _check_optimize_result(
[CV 4/5] END ..multi_class=ovr;, score=0.845 total time= 11.3s
/home/alton/.local/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(
[CV 5/5] END ..multi_class=ovr;; score=0.839 total time= 12.8s
/home/alton/.local/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(
[CV 1/5] END ..multi_class=multinomial;; score=0.840 total time= 16.8s
/home/alton/.local/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(
[CV 2/5] END ..multi_class=multinomial;; score=0.842 total time= 9.1s
/home/alton/.local/lib/python3.9/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(
[CV 3/5] END ..multi_class=multinomial;; score=0.840 total time= 11.6s

```
/home/alton/.local/lib/python3.9/site-  
packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  

```

```
[CV 4/5] END ..multi_class=multinomial;; score=0.838 total time= 10.5s
```

```
/home/alton/.local/lib/python3.9/site-  
packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  

```

```
[CV 5/5] END ..multi_class=multinomial;; score=0.832 total time= 11.0s
```

```
/home/alton/.local/lib/python3.9/site-  
packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(  

```

```
[CV 1/5] END ..multi_class=auto;; score=0.840 total time= 16.4s
```

```
/home/alton/.local/lib/python3.9/site-  
packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed  
to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

regression
  n_iter_i = _check_optimize_result(
[CV 2/5] END ..multi_class=auto;; score=0.842 total time= 14.3s
/home/alton/.local/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

  n_iter_i = _check_optimize_result(
[CV 3/5] END ..multi_class=auto;; score=0.840 total time= 11.4s
/home/alton/.local/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

  n_iter_i = _check_optimize_result(
[CV 4/5] END ..multi_class=auto;; score=0.838 total time= 12.4s
/home/alton/.local/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

  n_iter_i = _check_optimize_result(
[CV 5/5] END ..multi_class=auto;; score=0.832 total time= 12.7s
/home/alton/.local/lib/python3.9/site-
packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear_model.html#logistic-](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

regression

```
n_iter_i = _check_optimize_result(
```

```
[26]: Pipeline(steps=[('minmax', MinMaxScaler()),
                        ('lr',
                         GridSearchCV(estimator=LogisticRegression(),
                                       param_grid={'multi_class': ['ovr', 'multinomial',
                                                                    'auto']},
                                       verbose=3))])
```

```
[27]: pred = pipeline.predict(X_test)
```

```
[28]: from sklearn.metrics import classification_report, confusion_matrix,
      ↪ accuracy_score
```

```
[29]: print(confusion_matrix(y_test, pred))
```

```
[[14562   27    1   66   32]
 [  647  476   22  112  101]
 [    6    3  703   12    0]
 [ 1051   26    1  534    2]
 [  981   14    0   23  433]]
```

```
[30]: print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
1	0.84	0.99	0.91	14688
2	0.87	0.35	0.50	1358
3	0.97	0.97	0.97	724
4	0.71	0.33	0.45	1614
5	0.76	0.30	0.43	1451
accuracy			0.84	19835
macro avg	0.83	0.59	0.65	19835
weighted avg	0.83	0.84	0.81	19835

1.6 KNN

```
[31]: from sklearn.neighbors import KNeighborsClassifier
```

```
[32]: error_rate = []
      for i in range(1, 5000, 500):
```

```

knn = Pipeline(
    [("minmax", MinMaxScaler()), ("knn",
    ↪KNeighborsClassifier(n_neighbors=i))]
)
knn.fit(X_train_1, y_train_1)
pred_i = knn.predict(X_val)
print(f"Neighbours {i}")
print(f"Loss: {np.mean(pred_i != y_val)}")
# Take the mean where prediction is not equal to actual
error_rate.append(np.mean(pred_i != y_val))

```

```

Neighbours 1
Loss: 0.004726493571968742
Neighbours 501
Loss: 0.053062767834635746
Neighbours 1001
Loss: 0.08343836652382153
Neighbours 1501
Loss: 0.11828837912780439
Neighbours 2001
Loss: 0.14866397781699017
Neighbours 2501
Loss: 0.1872321653642551
Neighbours 3001
Loss: 0.2096672548525334
Neighbours 3501
Loss: 0.21785984371061257
Neighbours 4001
Loss: 0.22384673556843962
Neighbours 4501
Loss: 0.22920342828333753

```

```


[33]: plt.figure(figsize=(10, 6))
plt.plot(
    range(1, 5000, 500),
    error_rate,
    color="blue",
    linestyle="dashed",
    marker="o",
    markerfacecolor="red",
    markersize=10,
)
plt.title("Error Rate vs K Value")
plt.xlabel("K")
plt.ylabel("Error Rate")

```

```

[33]: Text(0, 0.5, 'Error Rate')

```



```
[34]: knn = Pipeline(
      [("minmax", MinMaxScaler()), ("knn", KNeighborsClassifier(n_neighbors=5))]
    )
      knn.fit(X_train_1, y_train_1)
```

```
[34]: Pipeline(steps=[('minmax', MinMaxScaler()), ('knn', KNeighborsClassifier())])
```

```
[35]: pred = knn.predict(X_test)
```

```
[36]: print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
1	1.00	1.00	1.00	14688
2	0.98	0.95	0.96	1358
3	1.00	1.00	1.00	724
4	0.99	0.99	0.99	1614
5	0.99	1.00	0.99	1451
accuracy			0.99	19835
macro avg	0.99	0.99	0.99	19835
weighted avg	0.99	0.99	0.99	19835


```
[37]: print(confusion_matrix(y_test, pred))
```

```
[[14659    20     2     3     4]
 [   47  1288     1    18     4]
 [    0     0   723     1     0]
 [    7     4     0  1603     0]
 [    6     0     0     1  1444]]
```

```
[38]: print(accuracy_score(y_test, pred))
```

```
0.9940509200907487
```

1.7 Decision Tree

```
[39]: from sklearn.tree import DecisionTreeClassifier
```

```
[40]: dt_params = {
        "criterion": ["gini", "entropy", "log_loss"],
        "splitter": ["best", "random"],
    }
    dt = Pipeline(
        [
            ("minmax", MinMaxScaler()),
            ("Bayes", GridSearchCV(DecisionTreeClassifier(), dt_params, verbose=3)),
        ]
    )
```

```
[41]: dt.fit(X_train_1, y_train_1)
```

Fitting 5 folds for each of 6 candidates, totalling 30 fits

```
[CV 1/5] END ...criterion=gini, splitter=best;; score=0.979 total time= 45.1s
[CV 2/5] END ...criterion=gini, splitter=best;; score=0.978 total time= 42.1s
[CV 3/5] END ...criterion=gini, splitter=best;; score=0.979 total time= 39.7s
[CV 4/5] END ...criterion=gini, splitter=best;; score=0.978 total time= 43.6s
[CV 5/5] END ...criterion=gini, splitter=best;; score=0.979 total time= 40.9s
[CV 1/5] END ...criterion=gini, splitter=random;; score=0.982 total time= 3.7s
[CV 2/5] END ...criterion=gini, splitter=random;; score=0.980 total time= 4.4s
[CV 3/5] END ...criterion=gini, splitter=random;; score=0.981 total time= 2.9s
[CV 4/5] END ...criterion=gini, splitter=random;; score=0.983 total time= 3.5s
[CV 5/5] END ...criterion=gini, splitter=random;; score=0.984 total time= 3.1s
[CV 1/5] END ..criterion=entropy, splitter=best;; score=0.984 total time= 34.7s
[CV 2/5] END ..criterion=entropy, splitter=best;; score=0.981 total time= 34.1s
[CV 3/5] END ..criterion=entropy, splitter=best;; score=0.983 total time= 33.4s
[CV 4/5] END ..criterion=entropy, splitter=best;; score=0.983 total time= 33.9s
[CV 5/5] END ..criterion=entropy, splitter=best;; score=0.981 total time= 33.3s
[CV 1/5] END criterion=entropy, splitter=random;; score=0.984 total time= 2.0s
[CV 2/5] END criterion=entropy, splitter=random;; score=0.983 total time= 2.3s
[CV 3/5] END criterion=entropy, splitter=random;; score=0.984 total time= 2.3s
```

```

[CV 4/5] END criterion=entropy, splitter=random;; score=0.983 total time= 1.9s
[CV 5/5] END criterion=entropy, splitter=random;; score=0.983 total time= 2.1s
[CV 1/5] END .criterion=log_loss, splitter=best;; score=0.983 total time= 34.7s
[CV 2/5] END .criterion=log_loss, splitter=best;; score=0.982 total time= 33.6s
[CV 3/5] END .criterion=log_loss, splitter=best;; score=0.984 total time= 33.3s
[CV 4/5] END .criterion=log_loss, splitter=best;; score=0.983 total time= 35.3s
[CV 5/5] END .criterion=log_loss, splitter=best;; score=0.981 total time= 38.4s
[CV 1/5] END criterion=log_loss, splitter=random;; score=0.984 total time=
2.4s
[CV 2/5] END criterion=log_loss, splitter=random;; score=0.982 total time=
2.6s
[CV 3/5] END criterion=log_loss, splitter=random;; score=0.984 total time=
2.5s
[CV 4/5] END criterion=log_loss, splitter=random;; score=0.983 total time=
2.1s
[CV 5/5] END criterion=log_loss, splitter=random;; score=0.983 total time=
1.9s

```

```

[41]: Pipeline(steps=[('minmax', MinMaxScaler()),
                        ('Bayes',
                         GridSearchCV(estimator=DecisionTreeClassifier(),
                                       param_grid={'criterion': ['gini', 'entropy',
                                                                'log_loss'],
                                                  'splitter': ['best', 'random']},
                                       verbose=3))])

```

```

[42]: pred_val = dt.predict(X_val)

```

```

[43]: print(confusion_matrix(y_val, pred_val))

```

```

[[11663   40    1   24   22]
 [   52 1008    7   13    7]
 [    2    6  569    2    0]
 [   25   15    0 1250    1]
 [   20   10    0    2 1129]]

```

```

[44]: print(classification_report(y_val, pred_val))

```

	precision	recall	f1-score	support
1	0.99	0.99	0.99	11750
2	0.93	0.93	0.93	1087
3	0.99	0.98	0.98	579
4	0.97	0.97	0.97	1291
5	0.97	0.97	0.97	1161
accuracy			0.98	15868
macro avg	0.97	0.97	0.97	15868

weighted avg	0.98	0.98	0.98	15868
--------------	------	------	------	-------

```
[45]: pred = dt.predict(X_test)
```

```
[46]: print(confusion_matrix(y_test, pred))
```

```
[[14585    60     2    21    20]
 [   66  1255     5    16    16]
 [    4     4   715     1     0]
 [   22    17     0  1570     5]
 [   27     3     0     2  1419]]
```

```
[47]: print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
1	0.99	0.99	0.99	14688
2	0.94	0.92	0.93	1358
3	0.99	0.99	0.99	724
4	0.98	0.97	0.97	1614
5	0.97	0.98	0.97	1451
accuracy			0.99	19835
macro avg	0.97	0.97	0.97	19835
weighted avg	0.99	0.99	0.99	19835

1.8 Random Forest

```
[48]: from sklearn.ensemble import RandomForestClassifier
```

```
[49]: rand = Pipeline(
    [ ("minmax", MinMaxScaler()), ("rand",
    ↪ RandomForestClassifier(criterion="gini")) ]
)
```

```
[50]: rand.fit(X_train_1, y_train_1)
```

```
[50]: Pipeline(steps=[('minmax', MinMaxScaler()), ('rand', RandomForestClassifier())])
```

```
[51]: pred_val = rand.predict(X_val)
```

```
[52]: print(classification_report(y_val, pred_val))
```

	precision	recall	f1-score	support
1	0.99	1.00	1.00	11750

2	0.97	0.96	0.96	1087
3	1.00	0.99	0.99	579
4	0.99	0.97	0.98	1291
5	1.00	0.98	0.99	1161
accuracy			0.99	15868
macro avg	0.99	0.98	0.99	15868
weighted avg	0.99	0.99	0.99	15868

```
[53]: print(confusion_matrix(y_val, pred_val))
```

```
[[11731  18    0    1    0]
 [   38 1041    2    5    1]
 [    2    2  574    1    0]
 [   23   10    0 1258    0]
 [   24    1    0    0 1136]]
```

```
[54]: pred = rand.predict(X_test)
```

```
[55]: print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
1	0.99	1.00	1.00	14688
2	0.98	0.95	0.96	1358
3	1.00	1.00	1.00	724
4	0.99	0.98	0.99	1614
5	1.00	0.99	0.99	1451
accuracy			0.99	19835
macro avg	0.99	0.98	0.99	19835
weighted avg	0.99	0.99	0.99	19835

```
[56]: print(confusion_matrix(y_test, pred))
```

```
[[14667  15    1    3    2]
 [   59 1289    2    5    3]
 [    2    1  721    0    0]
 [   19   14    0 1581    0]
 [   16    1    0    2 1432]]
```

1.9 SVM

```
[57]: from sklearn.svm import SVC
```

```
[58]: svm_param_grid = {"C": [0.1, 1, 10, 100, 1000], "gamma": [1, 0.1, 0.01, 0.001, 0.0001]}
```

```
[59]: svm = Pipeline([("minmax", MinMaxScaler()), ("SVM", SVC(C=0.1, gamma=1))])
```

```
[ ]:
```

```
[60]: svm.fit(X_train_1, y_train_1)
```

```
[60]: Pipeline(steps=[('minmax', MinMaxScaler()), ('SVM', SVC(C=0.1, gamma=1))])
```

```
[61]: pred_val = svm.predict(X_val)
```

```
[62]: print(classification_report(y_val, pred_val))
```

	precision	recall	f1-score	support
1	0.98	0.99	0.99	11750
2	0.95	0.92	0.94	1087
3	0.99	0.99	0.99	579
4	0.97	0.93	0.95	1291
5	0.99	0.91	0.95	1161
accuracy			0.98	15868
macro avg	0.98	0.95	0.96	15868
weighted avg	0.98	0.98	0.98	15868

```
[63]: print(confusion_matrix(y_val, pred_val))
```

```
[[11687   34    0   21    8]
 [   73  996    5   11    2]
 [    1    3  574    1    0]
 [   83    9    0 1199    0]
 [   98    1    0    0 1062]]
```

```
[64]: pred = svm.predict(X_test)
```

```
[65]: print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
1	0.98	0.99	0.99	14688
2	0.95	0.90	0.93	1358
3	1.00	1.00	1.00	724
4	0.97	0.93	0.95	1614
5	0.98	0.92	0.95	1451

accuracy			0.98	19835
macro avg	0.98	0.95	0.96	19835
weighted avg	0.98	0.98	0.98	19835

```
[66]: print(confusion_matrix(y_test, pred))
```

```
[[14608   39    0   32    9]
 [  102 1228    2   14   12]
 [    1    1  722    0    0]
 [   95   17    0 1502    0]
 [   118    2    0    1 1330]]
```

1.10 Gaussian Naive Bayes

```
[67]: from sklearn.naive_bayes import GaussianNB
```

```
[68]: gnb = Pipeline([("minmax", MinMaxScaler()), ("GNB", GaussianNB())])
```

```
[69]: gnb.fit(X_train_1, y_train_1)
```

```
[69]: Pipeline(steps=[('minmax', MinMaxScaler()), ('GNB', GaussianNB())])
```

```
[70]: pred_val = gnb.predict(X_val)
```

```
[71]: print(classification_report(y_val, pred_val))
```

	precision	recall	f1-score	support
1	0.94	0.69	0.79	11750
2	0.44	0.56	0.50	1087
3	0.63	0.93	0.75	579
4	0.29	0.79	0.42	1291
5	0.54	0.72	0.62	1161

accuracy			0.70	15868
macro avg	0.57	0.74	0.62	15868
weighted avg	0.82	0.70	0.73	15868

```
[72]: print(confusion_matrix(y_val, pred_val))
```

```
[[8065  649  143 2232  661]
 [  23  612  173  236   43]
 [   0   39  540    0    0]
 [ 188   71    5 1021    6]
 [ 265    8    0   56  832]]
```

```
[73]: pred = gnb.predict(X_test)
```

```
[74]: print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
1	0.95	0.70	0.81	14688
2	0.44	0.55	0.49	1358
3	0.67	0.94	0.78	724
4	0.30	0.80	0.43	1614
5	0.56	0.72	0.63	1451
accuracy			0.71	19835
macro avg	0.58	0.74	0.63	19835
weighted avg	0.82	0.71	0.74	19835

```
[75]: print(confusion_matrix(y_test, pred))
```

```
[[10325  807  149 2660  747]
 [   38  748  174  324   74]
 [    1   44  678    1    0]
 [  217   91    6 1291    9]
 [  341   13    0   53 1044]]
```

2 ANN

```
[76]: from tensorflow.keras.models import Sequential
```

```
2023-07-05 10:24:47.054506: I tensorflow/tsl/cuda/cudart_stub.cc:28] Could not
find cuda drivers on your machine, GPU will not be used.
2023-07-05 10:24:47.210747: I tensorflow/tsl/cuda/cudart_stub.cc:28] Could not
find cuda drivers on your machine, GPU will not be used.
2023-07-05 10:24:47.212063: I tensorflow/core/platform/cpu_feature_guard.cc:182]
This TensorFlow binary is optimized to use available CPU instructions in
performance-critical operations.
To enable the following instructions: AVX2 FMA, in other operations, rebuild
TensorFlow with the appropriate compiler flags.
2023-07-05 10:24:49.061011: W
tensorflow/compiler/tf2tensorrt/utils/py_utils.cc:38] TF-TRT Warning: Could not
find TensorRT
```

```
[77]: from tensorflow.keras.layers import Dense, Dropout
```

```
[78]: from tensorflow.keras.callbacks import EarlyStopping
```

```
[79]: from keras.utils import to_categorical
```

```
[80]: early_stop = EarlyStopping(monitor="val_loss", mode="min", verbose=1,  
    ↪patience=25)
```

```
[81]: model = Sequential()
```

```
[82]: model.add(Dense(300, activation="relu"))
```

```
[83]: model.add(Dropout(0.5))
```

```
[84]: model.add(Dense(150, activation="relu"))
```

```
[85]: model.add(Dropout(0.5))
```

```
[86]: model.add(Dense(75, activation="relu"))
```

```
[87]: model.add(Dropout(0.5))
```

```
[88]: model.add(Dense(35, activation="relu"))
```

```
[89]: model.add(Dropout(0.5))
```

```
[90]: model.add(Dense(15, activation="relu"))
```

```
[91]: model.add(Dropout(0.5))
```

```
[92]: model.add(Dense(5, activation="softmax"))
```

```
[93]: model.compile(loss="categorical_crossentropy", optimizer="adam")
```

```
[94]: model.fit(  
    x=X_train_1,  
    y=to_categorical(y_train_1)[:, 1:],  
    epochs=600,  
    validation_data=(X_val, to_categorical(y_val)[:, 1:]),  
    batch_size=32,  
    callbacks=[early_stop],  
)
```

Epoch 1/600

1984/1984 [=====] - 11s 5ms/step - loss: 0.6362 -
val_loss: 0.3861

Epoch 2/600

1984/1984 [=====] - 10s 5ms/step - loss: 0.4144 -
val_loss: 0.2850

Epoch 3/600

1984/1984 [=====] - 10s 5ms/step - loss: 0.3508 -
val_loss: 0.2124

Epoch 4/600
1984/1984 [=====] - 10s 5ms/step - loss: 0.3092 -
val_loss: 0.1773
Epoch 5/600
1984/1984 [=====] - 11s 5ms/step - loss: 0.2719 -
val_loss: 0.1573
Epoch 6/600
1984/1984 [=====] - 11s 5ms/step - loss: 0.2562 -
val_loss: 0.1316
Epoch 7/600
1984/1984 [=====] - 11s 6ms/step - loss: 0.2406 -
val_loss: 0.1300
Epoch 8/600
1984/1984 [=====] - 11s 5ms/step - loss: 0.2297 -
val_loss: 0.1213
Epoch 9/600
1984/1984 [=====] - 11s 6ms/step - loss: 0.2280 -
val_loss: 0.1042
Epoch 10/600
1984/1984 [=====] - 11s 6ms/step - loss: 0.2090 -
val_loss: 0.1170
Epoch 11/600
1984/1984 [=====] - 11s 6ms/step - loss: 0.2083 -
val_loss: 0.1070
Epoch 12/600
1984/1984 [=====] - 11s 6ms/step - loss: 0.2004 -
val_loss: 0.0979
Epoch 13/600
1984/1984 [=====] - 11s 6ms/step - loss: 0.1961 -
val_loss: 0.1079
Epoch 14/600
1984/1984 [=====] - 11s 6ms/step - loss: 0.1876 -
val_loss: 0.0866
Epoch 15/600
1984/1984 [=====] - 11s 6ms/step - loss: 0.1791 -
val_loss: 0.0859
Epoch 16/600
1984/1984 [=====] - 11s 6ms/step - loss: 0.1778 -
val_loss: 0.0813
Epoch 17/600
1984/1984 [=====] - 11s 6ms/step - loss: 0.1747 -
val_loss: 0.0818
Epoch 18/600
1984/1984 [=====] - 11s 6ms/step - loss: 0.1743 -
val_loss: 0.0816
Epoch 19/600
1984/1984 [=====] - 12s 6ms/step - loss: 0.1716 -
val_loss: 0.0844

Epoch 20/600
1984/1984 [=====] - 11s 6ms/step - loss: 0.1668 -
val_loss: 0.0889

Epoch 21/600
1984/1984 [=====] - 11s 6ms/step - loss: 0.1635 -
val_loss: 0.0842

Epoch 22/600
1984/1984 [=====] - 11s 6ms/step - loss: 0.1624 -
val_loss: 0.0920

Epoch 23/600
1984/1984 [=====] - 12s 6ms/step - loss: 0.1621 -
val_loss: 0.0810

Epoch 24/600
1984/1984 [=====] - 11s 6ms/step - loss: 0.1637 -
val_loss: 0.0872

Epoch 25/600
1984/1984 [=====] - 11s 6ms/step - loss: 0.1607 -
val_loss: 0.0878

Epoch 26/600
1984/1984 [=====] - 11s 5ms/step - loss: 0.1612 -
val_loss: 0.0821

Epoch 27/600
1984/1984 [=====] - 11s 5ms/step - loss: 0.1553 -
val_loss: 0.0785

Epoch 28/600
1984/1984 [=====] - 10s 5ms/step - loss: 0.1531 -
val_loss: 0.0749

Epoch 29/600
1984/1984 [=====] - 11s 5ms/step - loss: 0.1549 -
val_loss: 0.0856

Epoch 30/600
1984/1984 [=====] - 10s 5ms/step - loss: 0.1564 -
val_loss: 0.0721

Epoch 31/600
1984/1984 [=====] - 11s 5ms/step - loss: 0.1503 -
val_loss: 0.0797

Epoch 32/600
1984/1984 [=====] - 10s 5ms/step - loss: 0.1552 -
val_loss: 0.0732

Epoch 33/600
1984/1984 [=====] - 11s 5ms/step - loss: 0.1500 -
val_loss: 0.0739

Epoch 34/600
1984/1984 [=====] - 11s 6ms/step - loss: 0.1444 -
val_loss: 0.0682

Epoch 35/600
1984/1984 [=====] - 10s 5ms/step - loss: 0.1547 -
val_loss: 0.0729

Epoch 36/600
1984/1984 [=====] - 11s 6ms/step - loss: 0.1491 -
val_loss: 0.0668
Epoch 37/600
1984/1984 [=====] - 11s 5ms/step - loss: 0.1497 -
val_loss: 0.0641
Epoch 38/600
1984/1984 [=====] - 10s 5ms/step - loss: 0.1439 -
val_loss: 0.0710
Epoch 39/600
1984/1984 [=====] - 11s 5ms/step - loss: 0.1506 -
val_loss: 0.0684
Epoch 40/600
1984/1984 [=====] - 10s 5ms/step - loss: 0.1373 -
val_loss: 0.0713
Epoch 41/600
1984/1984 [=====] - 11s 6ms/step - loss: 0.1434 -
val_loss: 0.0677
Epoch 42/600
1984/1984 [=====] - 10s 5ms/step - loss: 0.1417 -
val_loss: 0.0767
Epoch 43/600
1984/1984 [=====] - 10s 5ms/step - loss: 0.1415 -
val_loss: 0.0708
Epoch 44/600
1984/1984 [=====] - 11s 5ms/step - loss: 0.1507 -
val_loss: 0.0650
Epoch 45/600
1984/1984 [=====] - 11s 5ms/step - loss: 0.1439 -
val_loss: 0.0724
Epoch 46/600
1984/1984 [=====] - 10s 5ms/step - loss: 0.1487 -
val_loss: 0.0691
Epoch 47/600
1984/1984 [=====] - 11s 5ms/step - loss: 0.1398 -
val_loss: 0.0787
Epoch 48/600
1984/1984 [=====] - 11s 5ms/step - loss: 0.1477 -
val_loss: 0.0849
Epoch 49/600
1984/1984 [=====] - 10s 5ms/step - loss: 0.1405 -
val_loss: 0.0660
Epoch 50/600
1984/1984 [=====] - 11s 5ms/step - loss: 0.1400 -
val_loss: 0.0718
Epoch 51/600
1984/1984 [=====] - 11s 5ms/step - loss: 0.1519 -
val_loss: 0.0691

```

Epoch 52/600
1984/1984 [=====] - 11s 5ms/step - loss: 0.1337 -
val_loss: 0.0689
Epoch 53/600
1984/1984 [=====] - 384s 194ms/step - loss: 0.1342 -
val_loss: 0.0746
Epoch 54/600
1984/1984 [=====] - 9s 4ms/step - loss: 0.1354 -
val_loss: 0.0782
Epoch 55/600
1984/1984 [=====] - 8s 4ms/step - loss: 0.1369 -
val_loss: 0.0796
Epoch 56/600
1984/1984 [=====] - 8s 4ms/step - loss: 0.1476 -
val_loss: 0.0704
Epoch 57/600
1984/1984 [=====] - 8s 4ms/step - loss: 0.1456 -
val_loss: 0.0695
Epoch 58/600
1984/1984 [=====] - 8s 4ms/step - loss: 0.1412 -
val_loss: 0.0807
Epoch 59/600
1984/1984 [=====] - 8s 4ms/step - loss: 0.1359 -
val_loss: 0.0677
Epoch 60/600
1984/1984 [=====] - 8s 4ms/step - loss: 0.1393 -
val_loss: 0.0711
Epoch 61/600
1984/1984 [=====] - 9s 4ms/step - loss: 0.1427 -
val_loss: 0.0687
Epoch 62/600
1984/1984 [=====] - 9s 5ms/step - loss: 0.1326 -
val_loss: 0.0694
Epoch 62: early stopping


```

```
[94]: <keras.callbacks.History at 0x7fce885ebe20>
```

```
[95]: losses = pd.DataFrame(model.history.history)
```

```
[96]: losses.plot()
```

```
[96]: <Axes: >
```



```
[97]: probs = model.predict(X_test)
```

```
620/620 [=====] - 1s 2ms/step
```

```
[98]: preds = np.argmax(probs, axis=1)
```

```
[99]: preds = preds + 1
```

```
[100]: y_test
```

```
[100]: array([3, 1, 1, ..., 1, 2, 1])
```

```
[101]: print(classification_report(y_test, preds))
```

	precision	recall	f1-score	support
1	0.99	1.00	0.99	14688
2	0.98	0.93	0.95	1358
3	1.00	1.00	1.00	724
4	0.98	0.99	0.98	1614
5	0.99	0.99	0.99	1451
accuracy			0.99	19835

macro avg	0.99	0.98	0.98	19835
weighted avg	0.99	0.99	0.99	19835

```
[102]: print(confusion_matrix(y_test, preds))
```

```
[[14638    21     1    18    10]
 [   84  1257     1    11     5]
 [    1     0   722     1     0]
 [   10     9     0  1592     3]
 [   10     0     0     1  1440]]
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[ ]:
```